

# Tittel

Forfatter

## 1 Backpropagation through time - BPTT

Definitions:

$$\mathbf{a}^{(t)} = \mathbf{b} + \mathbf{W}\mathbf{h}^{(t-1)} + \mathbf{U}\mathbf{x}^{(t)} \quad (1)$$

$$\mathbf{h}^{(t)} = \text{activation}(\mathbf{a}^{(t)}) \quad (2)$$

$$\mathbf{o}^{(t)} = \mathbf{c} + \mathbf{V}\mathbf{h}^{(t)} \quad (3)$$

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{o}^{(t)}) \quad (4)$$

$$\mathbf{U} = \text{input weights} \quad (5)$$

$$\mathbf{W} = \text{hidden weights} \quad (6)$$

$$\mathbf{V} = \text{output weights} \quad (7)$$

The goal is to maximize the probability of the observed data by estimating parameters (here:  $\mathbf{U}, \mathbf{V}, \mathbf{W}, \mathbf{b}, \mathbf{c}$ ). The estimated parameters yielding the highest maximum likelihood are called the maximum likelihood estimates. These parameters can be estimated by minimizing the cross-entropy between the model distribution and the data. This cross-entropy function, which is a function of inputs ( $\mathbf{x}$ ) and outputs ( $\mathbf{y}$ ), can be seen in (8).

$$C(\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(\tau)}\}, \{\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(\tau)}\}) \quad (8)$$

$$= \sum_t L^{(t)} \quad (9)$$

$$= - \sum_t \log p(y^{(t)} | \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(t)}\}) \quad (10)$$

$$(11)$$

Note:  $y^{(t)}$  is an entry in the output vector  $\hat{\mathbf{y}}^{(t)}$

The cost function is the negative log-likelihood function. Minimising this function is the same as maximum the likelihood of the parameters - not due to the log, but due to the negative prefix. The log in log-likelihood works, but I don't know why it is used?

Below we derive the gradients of the nodes in the computational graph from the deep learning book. These gradients must propagate backwards through time, from time  $t = \tau$  down to  $t = 0$ .

The gradient of the cost function at the output,  $\mathbf{o}$ , at time  $t$  is

$$\nabla_{\mathbf{o}^{(t)}} C = \frac{\delta C}{\delta \mathbf{o}^{(t)}} = \frac{\delta C}{\delta \mathbf{C}^{(t)}} \frac{\delta \mathbf{C}^{(t)}}{\delta \mathbf{o}^{(t)}} = \hat{\mathbf{y}}^{(t)} - \mathbf{1}_{i=y^{(t)}} \quad (12)$$

We have found a general expression for the gradient at the  $\mathbf{o}^{(t)}$ -nodes. The next step is to find an expression for the gradient of the final hidden (computational) node, at time  $\tau$ :  $\mathbf{h}^{(\tau)}$ . Its only descendant is  $\mathbf{o}^{(\tau)}$ , which means its gradient is solely dependent on this node, which makes it a good starting point for the later gradient calculations.

$$\nabla_{\mathbf{h}^{(\tau)}} C = (\nabla_{\mathbf{o}^{(\tau)}} C) \frac{\delta \mathbf{o}^{(\tau)}}{\delta \mathbf{h}^{(\tau)}} \quad (13)$$

$$= (\nabla_{\mathbf{o}^{(\tau)}} C) \mathbf{V} \quad (14)$$

$$\nabla_{\mathbf{h}^{(\tau)}} C = \mathbf{V}^\top (\nabla_{\mathbf{o}^{(\tau)}} C) \quad (15)$$

Where all the right hand side terms are known from before.

The only nodes that need gradient computation now, are all the hidden states preceding the last. I.e., for  $\mathbf{h}^{(t)}$ , where  $t = \{0, \dots, \tau - 1\}$ . For these time steps, the gradient is influenced by both the gradient at  $\mathbf{o}^{(t)}$ , as well as all the preceding hidden state gradients. Remember that the preceding hidden state of  $\mathbf{h}^{(t)}$  is  $\mathbf{h}^{(t+1)}$ , which has preceding hidden state  $\mathbf{h}^{(t+2)}$ , and so on. We are calculating the gradient starting from  $t = \tau - 1$ , working our way down to  $t = 0$ :

$$\nabla_{\mathbf{h}^{(\tau-1)}} C = \nabla_{\mathbf{o}^{(\tau-1)}} C \frac{\delta \mathbf{o}^{(\tau-1)}}{\delta \mathbf{h}^{(\tau-1)}} + \nabla_{\mathbf{h}^{(\tau)}} C \frac{\delta \mathbf{h}^{(\tau)}}{\delta \mathbf{h}^{(\tau-1)}} \quad (16)$$

$$\nabla_{\mathbf{h}^{(\tau-2)}} C = \nabla_{\mathbf{o}^{(\tau-2)}} C \frac{\delta \mathbf{o}^{(\tau-2)}}{\delta \mathbf{h}^{(\tau-2)}} + \nabla_{\mathbf{h}^{(\tau-1)}} C \frac{\delta \mathbf{h}^{(\tau-1)}}{\delta \mathbf{h}^{(\tau-2)}} \quad (17)$$

$$\nabla_{\mathbf{h}^{(\tau-3)}} C = \nabla_{\mathbf{o}^{(\tau-3)}} C \frac{\delta \mathbf{o}^{(\tau-3)}}{\delta \mathbf{h}^{(\tau-3)}} + \nabla_{\mathbf{h}^{(\tau-2)}} C \frac{\delta \mathbf{h}^{(\tau-2)}}{\delta \mathbf{h}^{(\tau-3)}} \quad (18)$$

$$= \nabla_{\mathbf{o}^{(\tau-3)}} C \frac{\delta \mathbf{o}^{(\tau-3)}}{\delta \mathbf{h}^{(\tau-3)}} + \left( \nabla_{\mathbf{o}^{(\tau-2)}} C \frac{\delta \mathbf{o}^{(\tau-2)}}{\delta \mathbf{h}^{(\tau-2)}} + \nabla_{\mathbf{h}^{(\tau-1)}} C \frac{\delta \mathbf{h}^{(\tau-1)}}{\delta \mathbf{h}^{(\tau-2)}} \right) \frac{\delta \mathbf{h}^{(\tau-2)}}{\delta \mathbf{h}^{(\tau-3)}} \quad (19)$$

$$= \nabla_{\mathbf{o}^{(\tau-3)}} C \frac{\delta \mathbf{o}^{(\tau-3)}}{\delta \mathbf{h}^{(\tau-3)}} + \left( \nabla_{\mathbf{o}^{(\tau-2)}} C \frac{\delta \mathbf{o}^{(\tau-2)}}{\delta \mathbf{h}^{(\tau-2)}} + \left( \nabla_{\mathbf{o}^{(\tau-1)}} C \frac{\delta \mathbf{o}^{(\tau-1)}}{\delta \mathbf{h}^{(\tau-1)}} + \nabla_{\mathbf{h}^{(\tau)}} C \frac{\delta \mathbf{h}^{(\tau)}}{\delta \mathbf{h}^{(\tau-1)}} \right) \frac{\delta \mathbf{h}^{(\tau-1)}}{\delta \mathbf{h}^{(\tau-2)}} \right) \frac{\delta \mathbf{h}^{(\tau-2)}}{\delta \mathbf{h}^{(\tau-3)}} \quad (20)$$

Generally:

$$\nabla_{\mathbf{h}^{(t)}} C = \nabla_{\mathbf{o}^{(t)}} C \frac{\delta \mathbf{o}^{(t)}}{\delta \mathbf{h}^{(t)}} + \nabla_{\mathbf{h}^{(t+1)}} C \frac{\delta \mathbf{h}^{(t+1)}}{\delta \mathbf{h}^{(t)}} \quad (21)$$

Some parts of the equations above have been colored to emphasize the parts that we take with us from one gradient calculation to the next.

It is observable that the gradient at time  $t$  is indeed dependent on all later time steps  $t + 1, t + 2, t + n$ , as well as the current output. The influence from current output can be seen before the first (+), while the influence from previous states can be observed after the first (+).

The final step is to calculate the gradient of the parameter nodes  $\mathbf{U}, \mathbf{W}, \mathbf{b}, \mathbf{V}, \mathbf{c}$ . Calculating the derivative of  $\mathbf{h}^{(t)}$  involves differentiating the activation function using the chain rule. In the equations below, the chain rule is applied to differentiate the activations with respect to different parameters, but the activation function itself is not differentiated because it depends

on which activation function in use. It is instead denoted as  $\nabla_{activation}$ .

$$\begin{aligned}
\nabla_{\mathbf{c}} C &= \sum_t (\nabla_{\mathbf{o}^{(t)}} C) \frac{\delta \mathbf{o}^{(t)}}{\delta \mathbf{c}^{(t)}} &= \sum_t (\nabla_{\mathbf{o}^{(t)}} C) \cdot 1 \\
\nabla_{\mathbf{v}} C &= \sum_t (\nabla_{\mathbf{o}^{(t)}} C) \frac{\delta \mathbf{o}^{(t)}}{\delta \mathbf{h}^{(t)}} &= \sum_t (\nabla_{\mathbf{o}^{(t)}} C) \mathbf{h}^{(t)} \\
\nabla_{\mathbf{b}} C &= \sum_t (\nabla_{\mathbf{h}^{(t)}} C) \frac{\delta \mathbf{h}^{(t)}}{\delta \mathbf{b}^{(t)}} &= \sum_t (\nabla_{\mathbf{h}^{(t)}} C) (\nabla_{activation}) \cdot 1 \\
\nabla_{\mathbf{w}} C &= \sum_t (\nabla_{\mathbf{h}^{(t)}} C) \frac{\delta \mathbf{h}^{(t)}}{\delta \mathbf{W}^{(t)}} &= \sum_t (\nabla_{\mathbf{h}^{(t)}} C) (\nabla_{activation}) \cdot \mathbf{h}^{(t-1)} \\
\nabla_{\mathbf{U}} C &= \sum_t (\nabla_{\mathbf{h}^{(t)}} C) \frac{\delta \mathbf{h}^{(t)}}{\delta \mathbf{U}^{(t)}} &= \sum_t (\nabla_{\mathbf{h}^{(t)}} C) (\nabla_{activation}) \cdot \mathbf{x}^{(t-1)}
\end{aligned}$$

The code written for BPTT is tricky to read. Below is conversions from math to code for gradient calculation of eq. (21):

$$\begin{aligned}
\nabla_{\mathbf{o}^{(t)}} C &= \hat{y}^{(t)} - \mathbf{1}_{i=y^{(t)}} && \Rightarrow \text{w\_hy} \\
\frac{\delta \mathbf{o}^{(t)}}{\delta \mathbf{h}^{(t)}} &= \mathbf{V} && \Rightarrow \text{grad\_o\_Cost} \\
\nabla_{\mathbf{h}^{(t+1)}} C &= \nabla_{\mathbf{o}^{(t+1)}} C \frac{\delta \mathbf{o}^{(t+1)}}{\delta \mathbf{h}^{(t+1)}} + \nabla_{\mathbf{h}^{(\tau)}} C \frac{\delta \mathbf{h}^{(t+2)}}{\delta \mathbf{h}^{(t+1)}} && \Rightarrow \text{prev\_grad\_h\_Cost} \\
\frac{\delta \mathbf{h}^{(t+1)}}{\delta \mathbf{h}^{(t)}} &= \mathbf{W} && \Rightarrow \text{w\_hh} \\
\nabla_{\mathbf{h}^{(t)}} C &&& \Rightarrow \text{grad\_h\_Cost}
\end{aligned}$$